

MULTI-CLOUD AI FINOPS

Multi-Cloud AI Tokenomics & FinOps

CostOps Portal — User Guide

The CostOps portal is the governed control plane for multi-cloud AI FinOps. It unifies AI token telemetry and pricing from five clouds onto one Microsoft Fabric lakehouse, attributes 100% of spend, runs twelve predictive ML models, and answers plain-language FinOps questions through a Data Agent. This guide walks through every screen and how to use it.

Michael Yaacoub | Sr Solution Engineer

Version 1.0 • September 2026

Introduction & Navigation

The CostOps portal is the governed control plane for multi-cloud AI FinOps. It unifies AI token telemetry and pricing from five clouds onto one Microsoft Fabric lakehouse, attributes 100% of spend, runs twelve predictive ML models, and answers plain-language FinOps questions through a Data Agent. This guide walks through every screen and how to use it.

Navigation at a glance

Screen	Icon	What it is for
Overview	grid	Multi-cloud KPI command center and trends.
Allocation	people	Cost attribution and chargeback by team, project, BU.
Model Pricing	pricetag	Live rate cards and a cost calculator across providers.
Power BI Reports	bar-chart	Embedded executive and ML dashboards.
12 ML Insights	sparkles	Outputs of the twelve PySpark ML use cases.
Data Agent Chat	chatbubbles	Ask FinOps questions in natural language.
Showcase	trending-up	Business value, architecture, and try-it links.

1. Business Value Showcase

The screenshot displays the 'Multi-Cloud AI Tokenomics & FinOps' landing page. At the top, a dark blue header contains the title and a subtitle: 'Govern every AI token across five clouds and fourteen model families on one governed lakehouse — with live cost attribution, twelve predictive ML models, interactive reports, and a Data Agent that answers FinOps questions in plain language.' Below this, four key metrics are presented in white text on a dark blue background: '5 CLOUDS UNIFIED', '281B TOKENS / MONTH', '12 PYSPARK ML MODELS', and '100% ATTRIBUTED SPEND'. A 'VALUE PROPOSITION' section follows, featuring the headline 'Turn ungoverned multi-cloud AI spend into billed savings' and a description: 'Govern, attribute, and optimize every token — a quality-safe FinOps control plane that cuts 30–45% of avoidable AI cost.' To the right of this text, two large blue figures are shown: '\$1.05M NET SAVINGS / MONTH' and '\$12.65M ANNUAL NET RUN-RATE'. A navigation bar at the bottom includes tabs for 'Video series', 'Business value', 'Features', 'Architecture', and 'Try it'. The 'Video series' tab is active, showing a video player with 'Video 01' and a duration of '05:12'.

MULTI-CLOUD AI FINOPS

Multi-Cloud AI Tokenomics & FinOps

Govern every AI token across five clouds and fourteen model families on one governed lakehouse — with live cost attribution, twelve predictive ML models, interactive reports, and a Data Agent that answers FinOps questions in plain language.

5	281B	12	100%
CLOUDS UNIFIED	TOKENS / MONTH	PYSPARK ML MODELS	ATTRIBUTED SPEND

VALUE PROPOSITION

Turn ungoverned multi-cloud AI spend into billed savings

Govern, attribute, and optimize every token — a quality-safe FinOps control plane that cuts 30–45% of avoidable AI cost.

\$1.05M	\$12.65M
NET SAVINGS / MONTH	ANNUAL NET RUN-RATE

[Video series](#) [Business value](#) [Features](#) [Architecture](#) [Try it](#)

Video 01 05:12

Business Value Showcase

The Showcase is a standalone, unauthenticated landing experience that presents the platform's value proposition, business case, features, architecture, and try-it links. It is the best place to start a briefing.

How to use it

- Open the Showcase from the portal or directly at the /showcase route.
- Read the value-proposition gadget at the top, then select Explore the business value to drill into the model.
- Use the tabs — Video series, Business value, Features, Architecture, Try it — to move between sections.
- In Try it, open the user guide, setup guide, and executive business case, or draft a source-access email.

Key elements


- Value proposition gadget with headline KPIs and a drill-down button.
- Business value tab: spend scenarios (\$1M / \$5M / \$10M per month), savings levers, and auditable assumptions.
- Architecture tab: the end-to-end architecture diagram.

2. Secure Sign-In

Secure Sign-In

Screenshot placeholder

Replace this image with an actual capture of the screen.



1600 x 900 • placeholder

Secure Sign-In

Access to the governed dashboards requires enterprise sign-in. Identity follows the user (On-Behalf-Of), so every downstream query to Fabric and the Data Agent runs with the signed-in user's permissions.


How to use it

- Select Sign in and authenticate with your organizational (Entra ID) account.
- Approve any one-time consent prompt for the delegated Power BI and Fabric scopes.
- Once the banner shows Tenant Connected, the dashboards load live governed telemetry.
- Use Sign out (top-right) to end the session on a shared device.

Key elements

- Entra ID sign-in with delegated, least-privilege scopes.
- Tenant Connected status indicator and Sign out control.

3. Token Command Center (Overview)



Token Command Center (Overview)

The Overview is the multi-cloud KPI command center. It summarizes tokens, cost, savings, latency, and operational health across all five providers for the selected window.


How to use it

- Scan the KPI cards: Total Tokens, Billed Cost, Negotiated Savings, P95 Latency, and Cost per Request.
- Use the Observation window (24H / 7D / 30D / 90D) to change the time range.
- Use the Cloud / Provider filter to focus on All Clouds or a single provider (AWS, GCP, Azure, OpenAI, Claude).
- Review the volume and cost trend charts, provider market share, and operational health meters.

Key elements

- Multi-cloud KPI cards with negotiated-savings and latency.
- Volume and cost trend charts; provider market-share breakdown.
- Observation-window tabs and provider filter chips.

4. Cost Allocation & Chargeback



Allocation attributes every token and dollar to the organization: business unit, department, team, project, and cost center — chargeback-ready and reconciled to negotiated rates.

How to use it

- Choose the attribution dimension (team, project, department, or business unit).
- Read the percentage breakdown and budget tracking for each entity.
- Compare market cost against negotiated cost to quantify savings per group.
- Export or hand off the allocation view for chargeback and budget reviews.

Key elements

- Cross-team, project, department, and BU attribution tables.
- Percentage breakdown and budget tracking per entity.

5. Multi-Cloud Model Pricing

MULTI-CLOUD AI FINOPS CONTROL PLANE

Multi-Cloud Model Pricing & Rate Cards

Live Telemetry SIGN OUT

OBSERVATION WINDOW: 24H 7D 30D 90D

CLOUD / PROVIDER: All Clouds (5) AWS Bedrock GCP Vertex AI Azure OpenAI OpenAI Direct Anthropic Claude

UPDATED 3:56 PM

DYNAMIC TIME-SERIES RATE CARDS

Multi-Cloud Model Pricing Catalog

Synced with published pricing APIs and SHA-256 provenance hashes

Interactive FinOps Workload Cost Estimator

Simulate token volume costs across model providers with enterprise discount rate cards

Input Volume: 10M tokens Output Volume: 5M tokens

Search model name, family, or provider...

All Clouds (5) AWS Bedrock GCP Vertex AI Azure OpenAI OpenAI Direct Anthropic Claude

PROVIDER & MODEL	FAMILY / VER	INPUT (\$/1M)	CACHED (\$/1M)	OUTPUT (\$/1M)	CONTEXT	SIMULATED SPEND	DISCOUNT
AZURE OPENAI mai-flash	MAI Flash v1.0	\$0.25	\$0.06	\$1.00	0	\$6.6 \$7.5	-12%
GCP VERTEX AI Google Gemini-1.0	Google Gemini-1.0	\$0.20	\$0.08	\$0.50	0	\$13.64	-12%

Multi-Cloud Model Pricing

The Model Pricing catalog shows live, effective-dated rate cards across AWS Bedrock, GCP Vertex AI, Azure OpenAI, OpenAI Direct, and Anthropic Claude, with a built-in cost calculator.


How to use it

- Browse rate cards by provider and model family; each rate is SHA-256 provenanced and effective-dated.
- Compare input, cached-input, cache-write, and output token prices per million.
- Toggle market vs negotiated pricing to see enterprise-discount impact.
- Use the calculator to estimate cost for a given token volume and model.

Key elements

- Live rate cards for all five providers.
- Market vs negotiated comparison and a pricing calculator.

6. Power BI Reports & Dashboards



Power BI Reports & Dashboards

The Power BI screen embeds the governed executive and ML reports built directly on OneLake via DirectQuery — the same source of truth as the API and Data Agent, with no data copy.


How to use it

- Select a report or dashboard: Consumption & Cost Analytics, ML Insights, or the executive/ML dashboards.
- Apply the date/time, department, and provider slicers to focus the view.
- Drill through the multi-level matrices (User > Model Family > Provider > Model > Version).
- Use in-report export where enabled to share a snapshot.

Key elements

- Embedded DirectQuery reports and service dashboards.
- Date, department, and provider slicers with drill-through.

7. 12 ML Insights



This screen visualizes the outputs of the twelve PySpark ML use cases — from anomaly detection and cost forecasting to model optimization and quota management.


How to use it

- Select an ML use case to view its latest scored output.
- Review anomaly flags, budget-overrun risk probabilities, and prompt-quality scores.
- Inspect model-optimization downgrade savings and quota recommendations.
- Use the insights to prioritize quality-safe optimization actions.

Key elements

- Interactive tiles for all 12 ML use cases (ml.* outputs).
- Risk probabilities, savings estimates, and recommendations.

8. Data Agent Chat & Prompt Library



Tokenomics FinOps Analyst
Tokenomics AI Analyst → Governed Lakehouse (tokenomics) Governed

Tokenomics FinOps Analyst
Governed Lakehouse • Gold & ML Tables Connected ONLINE

03:56 PM

Summarize multi-cloud AI token spend and the top 3 models by cost over the last 30 days.

Last 30 days (since 2026-08-20): multi-cloud AI token spend (all providers)

- **Total tokens:** 49.26B (Input 38.20B, Output 11.07B) across 7.71M requests
- **Cost (USD):**
- **Market:** \$299,772.80
- **Negotiated:** \$263,800.06
- **Savings vs market:** \$35,972.74 (~12.0%)

Sources: ml.tokenomics_usage_fact via get_spend_summary(days=30)

Prompts Ask the Tokenomics Data Agent (e.g. 'Predict Q4 spend for Yield Engineering', 'Claude Opus vs GPT Astra cost'). Press Enter to send, Shift+Enter for a new line...

Data Agent Chat & Prompt Library

The Data Agent Chat turns the governed lakehouse into plain-language answers. It calls the Fabric Data Agent through the APIM gateway under the signed-in user's identity, backed by a configurable Saved Prompt Library.

How to use it

- Pick a saved prompt from a category, or type your own FinOps question.
- Read the structured answer, including SQL/DAX breakdowns and recommendations.
- Refine with follow-up questions; the agent keeps conversational context.
- Export or copy the conversation for sharing where enabled.

Key elements

- Conversational FinOps analyst over governed Delta tables.
- Categorized Saved Prompt Library (executive, anomalies, optimization, quality, quota).

Tips & Troubleshooting

Tips

- Every query runs under your identity (On-Behalf-Of); you only see data you are authorized to see.
- Provider and time-window filters on the Overview carry the context you use across the dashboards.
- Savings are quality-safe: approve optimizations, then reconcile realized savings against invoices.
- The Showcase Business value tab is illustrative planning — replace figures with your own actuals.

Troubleshooting

Symptom	What to check
Stuck on "Loading governed usage telemetry"	Confirm sign-in completed and the API/gateway is reachable from your network.
Power BI report does not embed	Verify Power BI embedding tenant settings and that your account has report access.
Data Agent returns no answer	Check Foundry -> APIM -> Fabric Data Agent connectivity and your delegated permissions.
Empty KPI cards	Widen the observation window or clear the provider filter; confirm the lakehouse has data.